

E-commerce Session Purchase Prediction



Erica Trofimov
Luana Borma Brugger
Matteo Marello
Yardh Vilgot Göran Berglund

Introduction

Context & Objective

- E-commerce growth increased the volume of digital behavioural data (clicks, views, cart actions, time-on-site).
- Analysing session behaviour helps improve User Experience and targeting, supporting data-driven decisions.
- Practical value: If we can predict purchases, we can spot what behaviours lead to buying and optimise the website around them.

Project objective:

- Build a predictive model that classifies whether an online session ends in a purchase using behavioural features observed during that session

Research Question & Data

Research Question:

- Can we predict whether an online session will end up in a purchase based on behavioural features during that session (e.g., events volume, product views, cart additions, time spent)?

Dataset:

- Predictors: event logs capturing view/ cart/ purchase actions and interaction intensity and timing.
- Dataset provides the label of purchase vs. non-purchase.
- Scope: Kaggle *Ecommerce Behavior Data from Multi-Category Store*, ~10GB, 7 months of data, study focused on November 2019.

Data Storage

Challenge: *Avoid massive downloads on disk while still ensuring accessibility*

STEP 1: initial situation

Location: **kaggle**

Size: **10 GB**

Format: **CSV**

Accessibility: **low**

Initially the dataset is stored on kaggle.com, it's in .csv format and is roughly **10GB**. It's not possible to **access** it directly from our local machine **without downloading** it on the disk.

STEP 2: conversion

Location: **kaggle**

Size: **2.5 GB**

Format: **parquet**

Accessibility: **low**

Conversion into parquet format **reduces size** by 75%, the new dataset contains the same information and it is **optimised** for analytical operations, being a column-based storage file.

STEP 3: upload to GSC

Location: 

Size: **2.5 GB**

Format: **parquet**

Accessibility: **high**

Uploading the dataset on Google Storage Cloud (GSC) **increases accessibility** because it allows to access the data from local machines **avoiding massive download**.

Data Cleaning

Challenge: *Huge dataset that cannot be loaded into R as a standard data frame to perform operations.*

Addressing the challenge

- Arrow for lazy evaluation
- Process without loading raw data into memory

Goal: Push collection as late as possible

Out-of-memory Operations

- Performing “heavy” operations to create a compact intermediate dataset before collection

Goal: Grouping, counting events, average price, session start and end time before collecting

Solution

In-memory Operations

- Some operations are not supported by Arrow’s lazy backend
- Time parsing and date logic to calculate time-related features (e.g. *is_weekend*) needs to be calculated in-memory

Goal: Drop all unnecessary columns and collect

Scalability

- If dataset grows, we again run into the same issue since we need to collect.

Goal: Leverage BigQuery to perform all time transformations

Trade-off: Orchestration complexity vs (unlimited) scalability

2.5 GB (GSC)



~ 500 MB(Estimated object size)

Data Analysis

Device

- **16 GB RAM**
- **CPU**-only execution (no GPU, no parallelization)

EDA


- ‘**Big N**’ Complexity
- Summary of the features
- Target binary variable: n_purchase

Visualizations

- Correlation Matrix → Multicollinearity
- **Barplot** of purchases → high imbalance (5,51% purchases)

Big Data Considerations

- Remove objects not needed
- Plot with a sample if the dataset is huge



Predictive Models

Challenge: Computing models locally without reaching the vector memory limit

Ridge Logistic Regression

Key Aspects

Used **biglasso**: memory-efficient package for Lasso and Ridge

Converted train data to **filebacked.big.matrix** to reduce memory consumption during training.

Memory usage and **training time** recorded to evaluate efficiency

Regular cleanup with **rm()** and **gc()**

Challenging Computations:

Fitting the iterative k-fold **Cross-validation** for tuning parameter λ in the local memory

XGBoost

Key Aspects

Fast C++ implementation in R using **xgboost** and **DALEX** packages


Handled unbalanced data via **scale_pos_weight** prevented overcomputation by **early stopping** after 20 rounds without AUC improvement

Scalability:

- Might need **cloud computing** (e.g. Spark) for Cross validation and run more optimal models on bigger datasets
- Storing the models will occupy more memory
 - Better to **plot** on smaller **samples**

Results

Top 10 features by total Gain (percent)




Key drivers (XGBoost)

- Cart activity dominates (~69% of Gain).
- In-session intent signals explain most of the predictive power.

Behavioural pattern:

- Purchase probability jumps from 0 to 1 cart event, then declines gradually with more cart events.
- One decisive cart action separates many buyers; multiple cart actions may reflect indecision.


Partial Dependence Plot for n_cart



Results

Model performance

- XGBoost AUC ~ 0.96
- Ridge AUC ~ 0.89



Confusion-matrix trade-off

- Ridge: low FP but misses most purchases → very low purchase recall / F1
- XGBoost: captures most purchases, but with more FP → higher recall / F1, better for identifying likely buyers.

Ridge Regression Confusion Matrix		
	0	1
0	5161789	279007
1	38915	30948

XGBoost Confusion Matrix		
	0	1
0	4714989	27931
1	485716	282024

Thank you!